

RDS Monitoring

We need to enable the Enhanced Monitoring on the RDS to enable AWS RDS metrics. Please follow the below-given steps to enable enhanced monitoring.

To use enhanced monitoring, we must create an IAM role and then enable Enhanced Monitoring.

To create an IAM role for Amazon RDS enhanced monitoring

1. Open the IAM console at <https://console.aws.amazon.com>.
2. In the navigation pane, choose Roles.
3. Choose **Create role**.
4. Choose the **AWS service** tab, and then choose **RDS** from the list of services.
5. Choose **RDS - Enhanced Monitoring**, and then choose **Next**.
6. Ensure that the **Permissions policies** shows **AmazonRDSEnhancedMonitoringRole**, and then choose **Next**.
7. For **Role name**, enter a name for your role. For example, enter **demo**.
8. The trusted entity for your role is the AWS service **monitoring.rds.amazonaws.com**.
9. Choose **Create role**.

To turn on Enhanced Monitoring

1. Scroll to **Additional configuration**.
2. In **Monitoring**, choose **Enable Enhanced Monitoring** for your DB instance or read replica.
3. Set the **Monitoring Role** property to the IAM role that you created in the previous step to permit Amazon RDS to communicate with Amazon CloudWatch Logs for you, or choose **Default** to have RDS create a role for you named rds-monitoring-role.
4. Set the **Granularity** property to the interval, in seconds, between points when metrics are collected for your DB instance or read replica. The **Granularity** property can be set to one of the following values: 1, 5, 10, 15, 30, or 60.

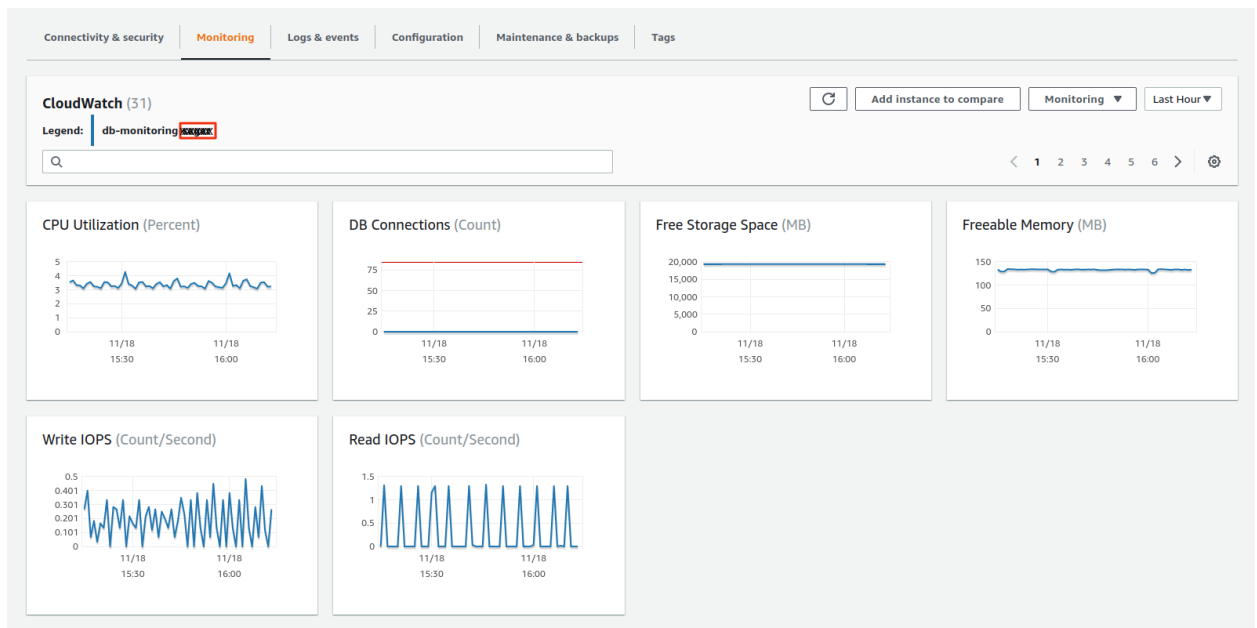
The fastest that the RDS console refreshes is every 5 seconds. If you set the granularity to 1 second in the RDS console, you still see updated metrics only every 5 seconds. You can retrieve 1-second metric updates by using CloudWatch Logs.

Cloudwatch to view metrics and for alerting

You can use cloudwatch to view metrics and implement alerts.

To view Enhanced metrics on cloudwatch

1. You can also view the metrics in the monitoring tab in the RDS console. This tab shows several metric graphs for each database.



2. You can use the cloudwatch console also to view the metrics. For that, Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
3. Under **Metrics** click on **All Metrics**.
4. Select the AWS region.
5. Click on **Browse** and select **RDS**, now you can select the option that how you want to see the RDS metrics.

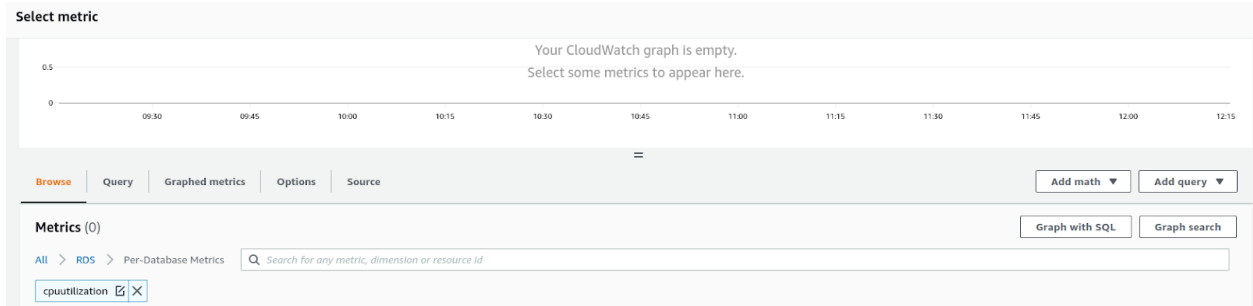
To implement alerts

Navigate to the cloudwatch console because the Cloudwatch alarms are created from the cloudwatch console.

Click on **create alarm**.



Click on **Select Metric** and type the name of the metric.



Choose the metric for the database you are going to monitor (which you can find in the tile labeled Per-Database Metrics). Click **Select Metric**.

Select metric

CPUUtilization

Browse Query Graphed metrics (1) Options Source

Add math Add query

Metrics (38)

Graph with SQL Graph search

All > RDS > Per-Database Metrics

Search for any metric, dimension or resource id

cpuutilization

	Metric name
	CPUUtilization
	CPUUtilization
<input checked="" type="checkbox"/>	db-monitoring-CPUUtilization
	CPUUtilization
	CPUUtilization
	CPUUtilization
	CPUUtilization
	CPUUtilization

Cancel Select metric

Select the threshold for your alarm.

Conditions

Threshold type



Static

Use a value as a threshold



Anomaly detection

Use a band as a threshold

Whenever WriteLatency is...

Define the alarm condition.



Greater

> threshold



Greater/Equal

>= threshold



Lower/Equal

<= threshold



Lower

< threshold

than...

Define the threshold value.

50

Must be a number

Click on **Next** and select the notification method for your alarm.

Configure actions

Notification

Alarm state trigger

Define the alarm state that will trigger this action.

Remove

☒ **In alarm**

The metric or expression is outside of the defined threshold.

☐ **OK**

The metric or expression is within the defined threshold.

☐ **Insufficient data**

The alarm has just started or not enough data is available.

Send a notification to the following SNS topic

Define the SNS (Simple Notification Service) topic that will receive the notification.

☐ Select an existing SNS topic

☒ **Create new topic**

☐ Use topic ARN to notify other accounts

Create a new topic...

The topic name must be unique.

Default_CloudWatch_Alarms_Topic

SNS topic names can contain only alphanumeric characters, hyphens (-) and underscores (_).

Email endpoints that will receive the notification...

Add a comma-separated list of email addresses. Each address will be added as a subscription to the topic above.

demo@demo.com

user1@example.com, user2@example.com

Create topic

Add notification

If your SNS topic is already created then you can select that and if you are creating a new topic then click on **Create topic**. Once done click on **Next**.

In the next window give your alarm a name.

Add name and description

Name and description

Alarm name

Alarm description - optional

Alarm description

Up to 1024 characters (0/1024)

Cancel

Previous

Next

Click on the **Next**. Now review your alarm and click on **Create alarm**.

CloudWatch > Alarms

Alarms (147) ☐ Hide Auto Scaling alarms

Clear selection

↺

Create c

×

Any state ▼

Any type ▼

Any actions ... ▼

<input type="checkbox"/>	Name	State	Last state update	Conditions
<input type="checkbox"/>	demo	⌚ Insufficient data	2022-11-18 17:56:03	WriteLatency >= 50 for 1 datapoints within 5 minutes

Now your alarm is created you will receive a notification whenever this alarm is triggered.

Viewing metrics in grafana.

To view cloudwatch metrics in grafana we have 2 options: -

1. Adding cloudwatch source in grafana.
2. Adding metrics from Prometheus.

Before moving further we need to create an IAM role with read access to cloudwatch so that grafana/ prometheus can read data from cloudwatch. Follow this [link](#) to create IAM role

1. Adding cloudwatch source.

In this approach, we will use cloudwatch as a data source and in the setting of the data source, you will get the option to add the connection details. You can select the SDK Default in which

you have to pass the Role ARN that you have created in the previous step. If you want to use an AWS IAM User to access the cloudwatch, then you have to pass the access key and secret key in the grafana or you can create a credentials file with the access key and secret key and upload it onto the grafana.

Connection Details

Authentication Provider	ⓘ	AWS SDK Default	🔍
Assume Role ARN	ⓘ	AWS SDK Default	
External ID	ⓘ	Access & secret key	
Endpoint	ⓘ	Credentials file	
Default Region	ⓘ	Choose	▼
Namespaces of Custom Metrics	ⓘ	Namespace1, Namespace2	

CloudWatch Logs

Once you have entered your authentication provider details, select the **Default Region** and click on **Save & test**. Now your cloudwatch data source is added.

Import the dashboard to access the metrics. You can use the dashboard ID **707**, to access the metrics. Your dashboard will look like this.



2. Adding metrics from Prometheus.

To add metrics from cloudwatch we need to create a yace exporter in our EKS cluster. For that, we need an IAM user who has the privilege to read data from cloudwatch.

Create a credentials file with the access key and secret key of the user. Once done run the below-mentioned command.

```
cat ./credentials | base64
```

Change the path of the credentials file according to your scenario. Now create a deployment file with the following manifest.

```
apiVersion: v1
kind: Namespace
metadata:
  name: yace
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: yace-rds
  namespace: yace
spec:
  selector:
    matchLabels:
      app: yace-rds
  replicas: 1
  template:
    metadata:
      labels:
        app: yace-rds
    annotations:
      prometheus.io/scrape: "true"
      prometheus.io/port: "5000"
  spec:
    containers:
      - name: yace
        image: quay.io/invisionag/yet-another-cloudwatch-exporter:v0.21.0-alpha
        ports:
          - containerPort: 5000
        volumeMounts:
          - name: yace-rds-config
```



```

        mountPath: /tmp/config.yml
        subPath: config.yml
      - name: yace-rds-credentials
        mountPath: /exporter/.aws/credentials
        subPath: credentials
    resources:
      limits:
        memory: "128Mi"
        cpu: "500m"
    volumes:
      - configMap:
          defaultMode: 420
          name: yace-rds-config
        name: yace-rds-config
      - secret:
          defaultMode: 420
          secretName: yace-rds-credentials
        name: yace-rds-credentials
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: yace-rds-config
  namespace: yace
data:
  config.yml: |
    discovery:
      jobs:
        - regions:
            - us-west-2
          type: rds
          enableMetricData: true
          metrics:
            - name: BinLogDiskUsage
              statistics:
                - Average
              period: 300
              length: 3600
            - name: BurstBalance
              statistics:
                - Average
              period: 300
              length: 3600
            - name: CPUUtilization
              statistics:

```

```
- Average
period: 300
length: 3600
- name: CPUCreditUsage
statistics:
  - Average
period: 300
length: 3600
- name: CPUCreditBalance
statistics:
  - Average
period: 300
length: 3600
- name: DatabaseConnections
statistics:
  - Average
period: 300
length: 3600
- name: DiskQueueDepth
statistics:
  - Average
  - Maximum
period: 300
length: 3600
- name: FailedSQLServerAgentJobsCount
statistics:
  - Average
period: 300
length: 3600
- name: FreeableMemory
statistics:
  - Average
period: 300
length: 3600
- name: FreeStorageSpace
statistics:
  - Average
period: 300
length: 3600
- name: MaximumUsedTransactionIDs
statistics:
  - Average
period: 300
length: 3600
- name: NetworkReceiveThroughput
```

```
    statistics:
      - Average
    period: 300
    length: 3600
  - name: NetworkTransmitThroughput
    statistics:
      - Average
    period: 300
    length: 3600
  - name: OldestReplicationSlotLag
    statistics:
      - Average
    period: 300
    length: 3600
  - name: ReadIOPS
    statistics:
      - Average
    period: 300
    length: 3600
  - name: ReadLatency
    statistics:
      - Maximum
      - Average
    period: 300
    length: 3600
  - name: ReadThroughput
    statistics:
      - Average
    period: 300
    length: 3600
  - name: ReplicaLag
    statistics:
      - Average
    period: 300
    length: 3600
  - name: ReplicationSlotDiskUsage
    statistics:
      - Average
    period: 300
    length: 3600
  - name: SwapUsage
    statistics:
      - Average
    period: 300
    length: 3600
```

```
- name: TransactionLogsDiskUsage
  statistics:
    - Average
    period: 300
    length: 3600
- name: TransactionLogsGeneration
  statistics:
    - Average
    period: 300
    length: 3600
- name: WriteIOPS
  statistics:
    - Average
    period: 300
    length: 3600
- name: WriteLatency
  statistics:
    - Maximum
    - Average
    period: 300
    length: 3600
- name: WriteThroughput
  statistics:
    - Average
    period: 300
    length: 3600
```

apiVersion: v1

kind: Secret

metadata:

name: yace-rds-credentials

namespace: yace

data:

Add in credentials the result of:

cat ./credentials | base64

credentials: |

XXXX

--

apiVersion: v1

kind: Service

metadata:

labels:

app: yace-rds

name: yace-svc

namespace: yace

```
spec:
  ports:
    - port: 5000
      protocol: TCP
      targetPort: 5000
  selector:
    app: yace-rds
```

In the above manifest replace the value of region with your region in configmap and in secret replace XXX with the value you have received by running command **cat ./credentials | base64** in the previous step.

Before moving further deploy Prometheus on the cluster with the following configuration.

Add the job exporter in your Prometheus.

```
- job_name: "yace-exporter"
  static_configs:
    - targets: ["yace-svc.yace.svc:5000"]
```

It will import the jobs from the yace pod and feed them to Prometheus. For alerting you need to add the below-mentioned alerting rule in Prometheus.

```
groups:
  - name: AWS-RDS
    rules:
      - alert: LongCPUThrottling
        expr: |
          aws_rds_cpuutilization_average > 95
        for: 10m
        labels:
          severity: page
        annotations:
          summary: CPU over 95% for 10 minutes in instance
            {{$labels.dimension_DBInstanceIdentifier}}
      - alert: LowFreeMemory
        expr: |
          aws_rds_freeable_memory_average < 128*1024*1024
        for: 10m
        labels:
          severity: page
        annotations:
          summary: Free memory under 128MB in instance
```

```

{{${labels.dimension_DBInstanceIdentifier}}}
- alert: LowFreeDisk
  expr: |
    aws_rds_free_storage_space_average < 512*1024*1024
  for: 10m
  labels:
    severity: page
  annotations:
    summary: Free disk under 512MB
{{${labels.dimension_DBInstanceIdentifier}}}
- alert: DiskFullIn48H
  expr: |
    predict_linear(aws_rds_free_storage_space_average[48h], 48 * 3600) < 0
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: Disk will be full within 48 hours in instance
{{${labels.dimension_DBInstanceIdentifier}}}
- alert: DiskFullIn12H
  expr: |
    predict_linear(aws_rds_free_storage_space_average[12h], 12 * 3600) < 0
  for: 10m
  labels:
    severity: page
  annotations:
    summary: Disk will be full within 12 hours in instance
{{${labels.dimension_DBInstanceIdentifier}}}
- alert: HighReadLatency
  expr: |
    aws_rds_read_latency_average > 0.250
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: Average read latency over 250ms in instance
{{${labels.dimension_DBInstanceIdentifier}}}
- alert: HighWriteLatency
  expr: |
    aws_rds_write_latency_average > 0.250
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: Average write latency over 250ms in instance

```

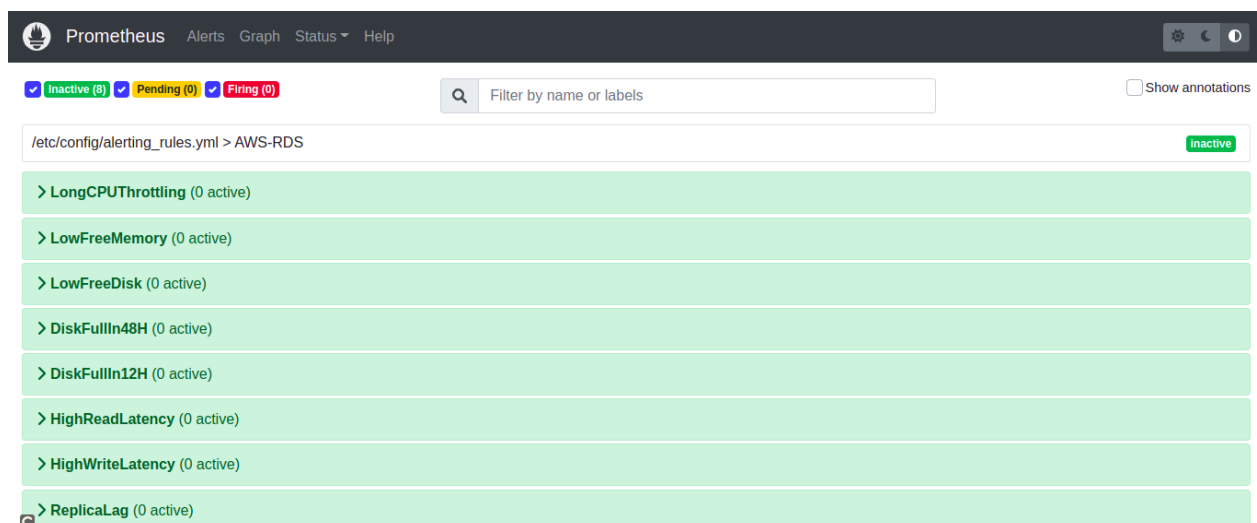
```

{{labels.dimension_DBInstanceIdentifier}}
- alert: HighDiskQueue
  expr: |
    aws_rds_disk_queue_depth_average > 25
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: High disk queue depth in instance
{{labels.dimension_DBInstanceIdentifier}}
- alert: ReplicaLag
  expr: |
    aws_rds_oldest_replication_slot_lag_average > 10
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: Average replication slot lag
{{labels.dimension_DBInstanceIdentifier}}

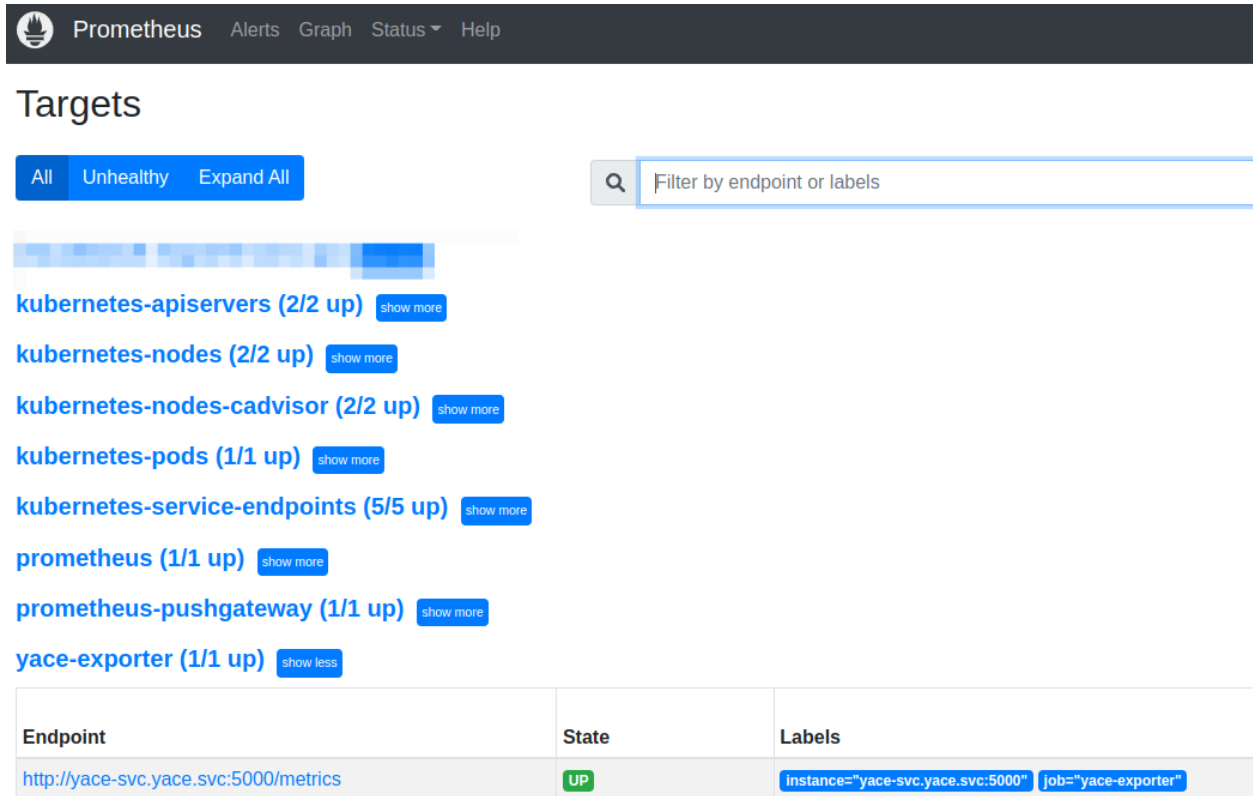
```

This configuration will enable the alerts on cloudwatch metrics, you can change the values in the configuration according to your need.

Once all the configuration is done you can see the alerts on the Prometheus.



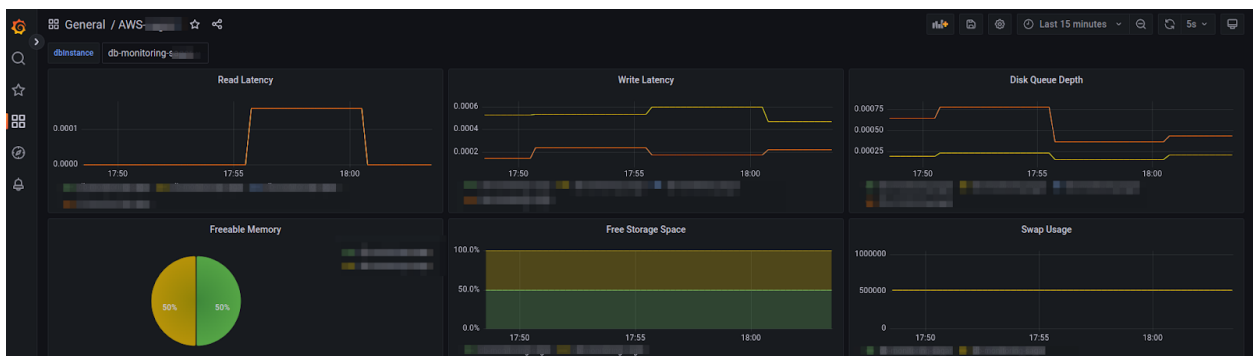
And under the targets, you can see the yace-exporter is up.



Now on grafana add Prometheus as a data source and import this dashboard. Use the wget command to download the dashboard.

```
wget
https://raw.githubusercontent.com/sagar-rafay/grafana-files/main/rds-promcat-dashboard.json
```

Once done import the dashboard on your grafana and explore the metrics. You can also implement alerts from the grafana dashboard.



References: -

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Monitoring.html

<https://promcat.io/apps/aws-rds>

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Grafana-support.html>

<https://www.bluematador.com/blog/how-to-monitor-amazon-rds-with-cloudwatch>

https://medium.com/@_oleksii_/using-aws-cloudwatch-in-grafana-8294b7a2e7dd